

# Statistical Comparison of PoS Committee Selection Algorithms

Robert Jones

23 April 2025

## 1 Simulation Experiment

We describe a Monte Carlo simulation experiment designed to evaluate two algorithms (A and B) that select committee members in a Proof-of-Stake (PoS) blockchain to achieve Byzantine Fault Tolerance (BFT) consensus. Their performance is measured by their robustness against faults when maintaining the 2/3 majority required for consensus. In this experiment, we sample a set of system configuration parameters  $\alpha$  from a joint distribution. For each sampled  $\alpha$ , we run many epochs (committee selections) for both algorithms and estimate a *fault tolerance score* (FTS) for both algorithms. The results are then analyzed using statistical tests to determine if there is a significant difference in FTS between the two algorithms.

For a given algorithm, let

$$p_f = \text{Probability that the system tolerates } f \text{ faults}$$

for  $f = 1, 2, 3, \dots$ . The **Fault Tolerance Score (FTS)** for the algorithm is then defined as:

$$\text{FTS} = \sum_{f=1}^{\infty} p_f,$$

which can be interpreted as the expected number of faults the system can tolerate before losing consensus.

- For **Algorithm A**, we denote the score as  $\text{FTS}_A$ .
- For **Algorithm B**, we denote the score as  $\text{FTS}_B$ .

A higher Fault Tolerance Score indicates that an algorithm can, on average, tolerate more faults while still preserving its 2/3 majority, thereby demonstrating better resilience.

To compare the algorithms, we examine the difference in their scores:

$$\Delta\text{FTS} = \text{FTS}_B - \text{FTS}_A.$$

- If  $\Delta\text{FTS} > 0$ , then **Algorithm B** is more resilient (i.e., it tolerates a higher number of faults on average).
- If  $\Delta\text{FTS} < 0$ , then **Algorithm A** performs better in terms of fault tolerance.
- If  $\Delta\text{FTS} \approx 0$ , then there is no significant difference between the two in terms of their fault tolerance capability.

## 2 Paired Statistical Testing

A crucial aspect of this experiment is the variation introduced by sampling  $\alpha$  from a distribution. By pairing the outcomes for A and B within each  $\alpha$ , we have structured the test as a randomized block design, where each block is a fixed environment  $\alpha$  and we compare two treatments (A vs B) within that block. This design improves sensitivity by removing inter-block (inter- $\alpha$ ) variability. Essentially, each scenario acts as its own control. If we instead ignored  $\alpha$  and pooled all epochs together, random differences in scenario difficulty would obscure the comparison.

Given that we have a sample of paired differences  $\Delta\text{FTS}(\alpha)$  for many random scenarios, the natural approach uses a paired statistical test to assess whether the mean of  $\Delta\text{FTS}(\alpha)$  is more significant than zero. The paired test treats the set of differences as a single sample. If the number of sampled scenarios is large enough, one common choice is the paired Student's t-test, which tests whether the mean difference is significantly different from zero. The paired t-test assumes the differences  $\Delta\text{FTS}(\alpha)$  are approximately normally distributed (an assumption often reasonable by the Central Limit Theorem if the number of samples is moderately large). It also assumes the sample of scenarios is independent. Here we know each  $\alpha$  is drawn independently, but we cannot assess normality in the data distribution.

### 2.1 Hypothesis Test

To evaluate whether Algorithm B is statistically superior to Algorithm A, we conduct a *one-sided hypothesis test*:

- **Null Hypothesis**,  $H_0$ : Algorithm B offers no improvement over A.

$$\text{mean}(\Delta\text{FTS}) = 0$$

- **Alternative Hypothesis**,  $H_1$ : Algorithm B is more fault tolerant.

$$\text{mean}(\Delta\text{FTS}) > 0$$

### 2.2 Bootstrap Sampling

Traditional statistical tests assume normality in the data distribution, which, after some testing, we determine is not the case here—rather, the data distribution is heavy tailed. Therefore, we will use a non-parametric alternative called *bootstrap sampling*. This approach is particularly useful when the underlying distribution is unknown or when the sample size is small. The bootstrap method allows us to estimate the sampling distribution of a statistic—in this case,  $\Delta\text{FTS}(\alpha)$  for a given parameter set  $\alpha$ —by resampling the observed data *with replacement*.

Now, we can conduct bootstrap resampling as follows:

1. **Resample**: Draw a large number (e.g., 10,000) of bootstrap samples from the centered differences (with replacement).
2. **Statistic**: For each bootstrap sample, compute its mean. This collection of sample means forms an empirical null distribution of the mean difference.
3. **Construct Empirical Null Distribution**: Shift the observed differences by subtracting the empirical mean (centered data) to simulate what would be expected under  $H_0$ .

4. **Compute Bootstrap Means:** For each bootstrap sample, calculate the mean difference in FTS scores.

5. **Formulate the One-Sided p-Value:** The one-sided p-value is calculated as

$$p = \frac{\text{number of bootstrap replicates with mean } \geq \text{observed mean}}{\text{total number of replicates}}.$$

6. **Interpret the p-value:** If this p-value is small (say, less than 0.05), it implies that such a high mean difference is very unlikely under the null hypothesis, lending support to the claim that algorithm B's performance is significantly better than algorithm A's.

7. **Conclusion:** If the p-value is small, we reject the null hypothesis  $H_0$  in favor of the alternative hypothesis  $H_1$ . This suggests that Algorithm B is statistically significantly better than Algorithm A.

### 3 Experimental Results

Appendix 1 contains the Python code used to conduct the Monte Carlo simulation experiment. The code implements the algorithms, generates the parameter configurations, and computes the Fault Tolerance Scores (FTS) for both Algorithm A and Algorithm B. It also includes the bootstrap sampling procedure to assess the statistical significance of the differences in FTS scores.

Running the code produces the results of the Monte Carlo simulation experiment, including the histograms of FTS scores for both algorithms and the histogram of the differences in FTS scores ( $\Delta$ FTS). The code also performs the bootstrap hypothesis test to determine whether Algorithm B outperforms Algorithm A.

The output of the code:

```
simulation trials: 100% [1000/1000] [13:54<00:00, 1.20it/s]
Observed Mean Difference (B - A): 1.733
One-sided p-value: 0.0000
Conclusion: Algorithm B is significantly better than Algorithm A.
```

#### 3.1 Discussion

The simulation trials were run for 1000 iterations, and the observed mean difference in Fault Tolerance Scores (FTS) between Algorithm B and Algorithm A was found to be 1.733. The one-sided p-value was less than 0.0001, indicating strong evidence against the null hypothesis  $H_0$ . Therefore, we conclude that Algorithm B is significantly better than Algorithm A in terms of fault tolerance.

The code also generates visualizations of the results—figures 1, 2, and 3—including histograms of the FTS scores for both algorithms and the histogram of the differences in FTS scores ( $\Delta$ FTS). The histograms provide a visual representation of the distribution of FTS values for both algorithms across multiple parameter configurations.

The figures illustrate the results of our Monte Carlo simulation experiment. Figure 1 shows the comparison of Fault Tolerance Scores (FTS) between Algorithm A and Algorithm B across different parameter configurations. The overlayed histograms in figure 2 displays the distribution of FTS




Figure 1: Comparison of Fault Tolerance Scores (FTS) between Algorithm A and Algorithm B across different parameter configurations. Higher scores indicate better fault tolerance.




Figure 2: Comparison of FTS histograms between Algorithm A and Algorithm B. The histogram shows the distribution of FTS values for both algorithms across multiple parameter configurations.




Figure 3: Histogram of the difference in FTS scores ( $\Delta\text{FTS}$ ) between Algorithm A and Algorithm B.

values for both algorithms. Finally, figure 3 shows the histogram of the difference in FTS scores ( $\Delta\text{FTS}$ ) between Algorithm A and Algorithm B. The mean of the  $\Delta\text{FTS}$  distribution is significantly greater than zero, with a p-value less than 0.05 from the bootstrap hypothesis test. This indicates that we can reject the null hypothesis  $H_0$  in favor of the alternative hypothesis  $H_1$ , indicating that Algorithm B consistently outperforms Algorithm A in terms of fault tolerance.

## 4 Conclusion

Our experimental results show the distribution of  $\Delta\text{FTS}$  values across multiple parameter configurations. As shown in figures, Algorithm B demonstrates consistently higher fault tolerance scores compared to Algorithm A. The bootstrap hypothesis test confirms that the observed differences are statistically significant, with a p-value less than 0.05. This suggests that Algorithm B is indeed more resilient in terms of fault tolerance, making it a preferable choice for PoS committee selection in Byzantine Fault Tolerance scenarios. The use of bootstrap sampling allowed us to overcome the limitations of traditional statistical tests, particularly in the presence of non-normal data distributions. By resampling the observed data and constructing an empirical null distribution, we were able to accurately assess the significance of the differences in FTS scores between the two algorithms. The results of this study provide valuable insights into the performance of PoS committee selection algorithms and highlight the importance of robust statistical methods in evaluating their effectiveness. Future work may involve further exploration of different parameter configurations, as well as the application of these methods to other PoS consensus algorithms.

## A Appendix: Code

```
1  #!/usr/bin/env python
2  # -*- coding: utf-8 -*-
3
4  import numpy as np
5  import pandas as pd
6  import scipy.stats as stats
7  import seaborn as sns
8  from matplotlib import pyplot as plt
9  from numpy import ceil, ndarray, ones, random, zeros
10 from tqdm import tqdm
11
12 from slot_alloc_sim import simulate_epoch_permissioned,
13     simulate_epoch_registered
14
15 def sample_multinomial(num_participants: int) -> ndarray:
16     """This function generates a probability vector from a flat Dirichlet
17     distribution
18     (with equal pseudo-counts for every participant)
19
20     Args:
21         num_participants: Number of participants (candidates from group)
22
23     Returns:
24         probabilities: A probability vector sampled from the Dirichlet
25     distribution
26     """
27
28     # Define a "flat" Dirichlet prior where every pseudo-count is 1.
29     alpha = ones(num_participants)
30
31
32     # Sample a probability vector from the Dirichlet distribution,
33     # which represents the multinomial probabilities for each participant.
34     return random.dirichlet(alpha)
35
36
37 def random_factor(n):
38     """
39     Computes all integer factors of a number and randomly selects one.
40
41     Args:
42         n (int): The number to find factors for
43
44     Returns:
45         int: A randomly selected factor of n
46     """
47
48     # Find all factors of n
49     factors = []
50     for i in range(1, int(n ** 0.5) + 1):
51         if n % i == 0:
52             factors.append(i)
53             if i != n // i: # Avoid adding the same factor twice for perfect
squares
54                 factors.append(n // i)
55
56     # If no factors are found (should only happen if n=0), return 1
57     if not factors:
58         return 1
```

```

54     # Return a random factor
55     return random.choice(factors)
56
57
58
59     # Define the function to sample alpha from a joint distribution
60     def sample_alpha() -> dict:
61         """
62             Generates a sample allocation of seats among participants using a
63             multinomial
64             distribution.
65
66             The function determines the total number of seats and the number of
67             participants from predefined options. It then generates probabilities for
68             each participant and uses these probabilities to allocate seats
69             proportionally. The result is returned as a dictionary containing all the
70             generated values.
71
72             Returns:
73                 dict: A dictionary containing the following keys and values:
74                     - total_seats (int): Total number of available seats.
75                     - num_participants (int): Number of participants.
76                     - p (List[float]): Probabilities for each participant.
77                     - seat_counts (List[int]): List containing the number of seats
78                         allocated to each participant.
79
80             Raises:
81                 AssertionError: If the internal logic fails, such as the sum of seat
82                 counts exceeding the specified total seats due to an unexpected error.
83             """
84
85             # Consider a random number of registered seats over this range
86             num_registered = random.choice([10, 20, 40, 80, 160, 320])
87
88             # Let num_permitted seats be a random percentage of num_registered
89             num_permitted = int(ceil(random.choice([0.1, 0.3, 0.5, 0.7]) *
90             num_registered))
91
92             # Choose a random factor for seats_per_federated
93             seats_per_federated = random_factor(num_permitted)
94
95             # This is guaranteed now to be an integer number of federated participants
96             num_federated = num_permitted // seats_per_federated
97
98             total_seats = num_registered + num_permitted
99
100            registered_probabilities = sample_multinomial(num_registered)
101
102            max_faults = 10
103
104            sample = dict(
105                total_seats=total_seats,
106                num_registered=num_registered,
107                num_permitted=num_permitted,
108                num_federated=num_federated,
109                seats_per_federated=seats_per_federated,
110                registered_probabilities=registered_probabilities,
111                max_faults=max_faults,
112                num_epochs=100,
113            )

```

```

111     return sample
112
113
114     def faults_tolerated(committee_seats: pd.Series) -> int:
115         """
116             Compute the number of faults tolerated by the committee
117
118         Args:
119             committee_seats (pd.Series): Series of committee seats.
120
121         Returns:
122             int: The number of faults tolerated by the committee.
123         """
124         voting_strength = committee_seats.sort_values(ascending=False).divide(
125             committee_seats.sum(),
126         )
127         threshold = 1 / 3 # BFT finality risk threshold
128         faults = np.where(np.cumsum(voting_strength) > threshold)[0][0]
129         return faults
130
131
132     def simulate_epoch_federated(
133         registered_seat_counts: ndarray,
134         num_federated: int,
135         seats_per_federated: int,
136         verbose: bool = False,
137     ) -> pd.DataFrame:
138         """
139             Creates a committee with registered and federated seats.
140
141             The 'simulate_epoch_federated' function allocates committee seats between
142             two types of nodes: registered nodes
143                 (whose seat allocation is proportional to their stake) and federated nodes
144                 (which receive a fixed number of
145                     seats). It calculates the voting strength of each group, ensures they sum
146                     to 1.0, and returns a structured
147                         DataFrame containing the committee composition with seat assignments for
148                         both node types.
149
150             Args:
151                 registered_seat_counts (ndarray): seat counts for registered
152                 participants.
153                 num_federated (int): Number of federated nodes.
154                 seats_per_federated (int): Number of seats per federated node.
155                 verbose (bool): Whether to print detailed information.
156                     Default is False.
157
158             Returns:
159                 pd.DataFrame: Committee seats information with kind and seat
160                 assignments.
161             """
162             # Cast the seat counts to a pandas Series
163             registered_seat_counts = pd.Series(registered_seat_counts, name=""
164             registered_seats")
165
166             # Get the number of registered participants (SPOs)
167             num_registered = registered_seat_counts.shape[0]
168
169             # Calculate the number of federated seats in the committee

```

```

163     federated_seats = seats_per_federated * num_federated
164
165     # Committee size is the total number of registered and federated seats
166     registered_seats = registered_seat_counts.sum()
167     committee_size = registered_seats + federated_seats
168
169     # Calculate the voting strength of the registered seats
170     registered_voting_strength = registered_seats / committee_size
171
172     # Calculate the voting strength of the federated seats
173     federated_voting_strength = federated_seats / committee_size
174
175     # Assert that the voting strengths sum to 1.0
176     assert (
177         registered_voting_strength + federated_voting_strength == 1.0
178     ), "Voting strength does not sum to 1.0"
179
180     # Create a series for the federated seats on the committee
181     federated_seat_counts = pd.Series(
182         ones(num_federated, dtype="int64") * seats_per_federated,
183         index=[str(i) for i in range(num_registered + 1, num_registered +
184             num_federated + 1)],
184         dtype="int64",
185         name="federated seats",
186     )
187
188     # Combine the federated and registered seats into a single DataFrame
189     committee_seats = (
190         pd.concat(
191             [federated_seat_counts, registered_seat_counts],
192             keys=["federated", "registered"],
193             names=["kind", "index"],
194             ignore_index=False,
195         )
196         .reset_index()
197         .rename(columns={0: "seats"})
198         .set_index("index")
199         .sort_values(by=["seats", "kind"], ascending=[False, True])
200     )
201     if verbose:
202         print(
203             f"Committee size .... = {committee_size}\n"
204             "-----\n"
205             "Registered:\n"
206             f"Number registered.. = {num_registered}\n"
207             f"Number of seats.... = {registered_seats}\n"
208             f"Voting strength.... = {registered_voting_strength:.2%}\n"
209             "-----\n"
210             "Federated:\n"
211             f"Number federated... = {num_federated}\n"
212             f"Seats per federated = {seats_per_federated}\n"
213             f"Number of seats.... = {federated_seats}\n"
214             f"Voting strength.... = {federated_voting_strength:.2%}\n",
215         )
216     return committee_seats
217
218
219     def simulate_committee_federated(
220         registered_probabilities: ndarray,

```

```

221     num_registered: int,
222     num_federated: int,
223     seats_per_federated: int,
224     num_epochs: int,
225     **kwargs,
226     ) -> pd.DataFrame:
227     """
228         Simulate the allocation of committee seats over multiple epochs.
229
230         This function simulates the process of assigning committee seats for a
231         specified
232         number of epochs using the simulate_epoch_federated function. It collects
233         seat
234         assignments across all epochs and returns them as a DataFrame.
235
236         Args:
237             registered_probabilities (ndarray): Probabilities for registered
238             participants.
239             num_registered (int): Number of registered participants.
240             num_federated (int): Number of federated nodes.
241             seats_per_federated (int): Number of seats per federated node.
242             num_epochs (int): Number of epochs to simulate.
243             kwargs: Additional keyword arguments.
244
245         Returns:
246             pd.DataFrame: DataFrame containing committee seat assignments for all
247             epochs,
248             with epochs as rows and committee members as columns.
249             """
250
251             committee_list = []
252             for epoch in range(num_epochs):
253                 # Sample the committee seat configuration from the multinomial
254                 # distribution
255                 # parameterized by registered_probabilities.
256                 registered_seat_counts = random.multinomial(num_registered,
257             registered_probabilities)
258
259                 # Assign committee seats for the current epoch
260                 committee = simulate_epoch_federated(
261                     registered_seat_counts=registered_seat_counts,
262                     num_federated=num_federated,
263                     seats_per_federated=seats_per_federated,
264                     )
265                 committee_list.append(committee.seats)
266
267             assert len(committee_list) == num_epochs, "Number of epochs does not match
268             ."
269
270             committee_seats = pd.concat(committee_list, keys=range(num_epochs), axis
271             =1).T
272
273             return committee_seats
274
275
276
277     def simulate_proposed(
278         registered_probabilities: ndarray,
279         num_registered: int,
280         num_permitted: int,
281         num_epochs: int,

```

```

272         **kwargs,
273     ) -> pd.DataFrame:
274         """Runs the proposed algorithm simulation. For each epoch, simulates both
275         the
276         registered and permissioned candidate selection. Stores the number of
277         slots
278         each candidate gets and returns the cumulative results.
279
280         Args:
281             registered_probabilities (ndarray): Probabilities for registered
282             candidates.
283             num_registered (int): Number of registered candidates.
284             num_permitted (int): Number of permitted (federated)nodes.
285             num_epochs (int): Number of epochs to simulate.
286             kwargs: Additional keyword arguments.
287
288         Returns:
289             pd.DataFrame: DataFrame containing committee seat assignments for all
290             epochs ,
291             with epochs as rows and committee members as columns.
292             """
293             permissioned_candidates = [f"P{i}" for i in range(num_permitted)]
294             permissioned_results = {name: [] for name in permissioned_candidates}
295             registered_results = {i: [] for i in range(num_registered)}
296
297             for _ in range(num_epochs):
298                 registered_candidates = {i: registered_probabilities[i] for i in
299                 registered_results.keys()}
300
301                 # Simulate this epoch
302                 reg_assign = simulate_epoch_registered(registered_candidates ,
303                 num_registered)
304                 perm_assign = simulate_epoch_permissioned(permissioned_candidates ,
305                 num_permitted)
306
307                 # Collect results
308                 for name in registered_candidates:
309                     registered_results[name].append(reg_assign[name])
310                     for name in permissioned_candidates:
311                         permissioned_results[name].append(perm_assign[name])
312
313                 # Combine the two dictionaries
314                 combined = registered_results.copy()
315                 combined.update(permissioned_results)
316
317                 # Convert to DataFrame
318                 committee_seats = (
319                     pd.DataFrame.from_dict(combined, orient="index")
320                     .astype(int)
321                     .fillna(0)
322                     .transpose()
323                 )
324                 return committee_seats
325
326
327             def calculate_fault_tolerance_probability(
328                 committee_seats: pd.DataFrame,
329                 fault_tolerance: int = 1,
330             ) -> float:

```

```

324     """
325     Calculate the probability of tolerating a given number of faults
326     in a committee of a given size.
327     The function uses Monte Carlo simulation to estimate the probability.
328     Args:
329         committee_seats (pd.DataFrame): DataFrame of committee seat
330             assignments of both registered and permissioned members
331             for each epoch.
332         fault_tolerance (int): Number of faults to tolerate.
333
334     Returns:
335         float: Estimated probability of tolerating the given number of faults.
336     """
337     if fault_tolerance == 0:
338         probability = 1.0
339     else:
340         # Calculate success rate through Monte Carlo simulation
341         probability = (
342             committee_seats.apply(faults_tolerated, axis=1) >= fault_tolerance
343         ).mean()
344     return probability
345
346
347     # Stub functions for Algorithm A and B
348     def algorithm(function, **params) -> float:
349         """
350             Computes fault tolerance probabilities using an algorithm that processes a
351             defined function and its
352             parameters to determine committee seats and calculate probabilities for a
353             fault tolerance metric.
354
355             Args:
356                 function (Callable): A function that determines committee seats based
357                     on provided parameters.
358                 **params: A set of keyword arguments to be passed into the function 'function'.
359                     Must include
360                         'max_faults' which defines the maximum faults for calculating
361                     probabilities; defaults to 5
362                         if not specified.
363
364             Returns:
365                 mean_fault_tolerance: (float) The mean fault tolerance probability for
366                     the algorithm.
367             """
368             committee_seats = function(**params)
369
370             # Calculate fault tolerance probabilities for the algorithm
371             faults = np.arange(1, params.get("max_faults", 5) + 1)
372             ftprob = pd.Series(0.0, index=faults, name="probability")
373             for f in faults:
374                 # p is the probability of tolerating f faults
375                 p = calculate_fault_tolerance_probability(committee_seats,
376 fault_tolerance=f)
377                 ftprob.loc[f] = p
378                 if p == 0: # since the rest of the series will be zero as well
379                     break
380
381             # Compute the performance score as the sum of probabilities
382             score = ftprob.sum()

```

```

376         return score
377
378
379
380     def hypothesis_test(
381         results_a: ndarray,
382         results_b: ndarray,
383         kind: str = "bootstrap",
384         n_bootstrap: int = 10000,
385     ):
386         """
387             Conducts a hypothesis test to compare two result sets, 'results_a' and 'results_b', based on the specified
388             kind of statistical testing approach. By default, a bootstrap approach is used to calculate the one-sided
389             p-value for the null hypothesis. Alternatively, a paired t-test can be performed. The results help determine
390             whether algorithm B (corresponding to 'results_b') is significantly better than algorithm A (corresponding
391             to 'results_a').
392
393             Args:
394                 results_a: An array of numerical results for algorithm A.
395                 results_b: An array of numerical results for algorithm B.
396                 kind: The kind of statistical test to perform. Options are "bootstrap" (default) or "paired_t".
397                 n_bootstrap: The number of bootstrap resampling iterations to perform.
398             Only applicable
399                 if the 'kind' is "bootstrap".
400
401             Raises:
402                 ValueError: If 'kind' is not "bootstrap" or "paired_t".
403                 RuntimeError: If inconsistencies in data prevent calculations, such as mismatched array lengths
404                         or insufficient data.
405             """
406
407         if kind == "bootstrap":
408             # Calculate the observed mean difference.
409             differences = results_b - results_a
410             obs_diff = np.mean(differences)
411
412             # For a one-sided test (B > A), if the observed mean difference is non
413             # positive,
414             # the p-value would be 1 (or you might decide not to proceed).
415             if obs_diff <= 0:
416                 p_value = 1.0
417             else:
418                 # Center the differences: adjust the data so that its mean becomes
419                 # 0
420                 # (the null hypothesis condition)
421                 differences_centered = differences - obs_diff
422
423                 # Bootstrap resampling from the centered differences.
424                 bootstrap_means = np.empty(n_bootstrap)
425                 for i in range(n_bootstrap):
426                     sample = np.random.choice(differences_centered, size=len(
427                     differences_centered), replace=True)
428                     bootstrap_means[i] = np.mean(sample)

```

```

425         # One-sided p-value: proportion of bootstrap means greater than or
426         # equal to the observed difference.
427         p_value = np.mean(bootstrap_means >= obs_diff)
428
429         print(f"Observed Mean Difference (B - A): {obs_diff:.3f}")
430         print(f"One-sided p-value: {p_value:.4f}")
431
432     else:
433         # Perform a paired t-test to determine if B is significantly better
434         # than A
435         t_stat, p_value = stats.ttest_rel(results_a, results_b, alternative='
436         less')
437
438         print(f"Mean mA: {results_a.mean():.3f}, Mean mB: {results_b.mean():.3
439         f}")
440         print(f"Paired t-test result: t-statistic = {t_stat:.3f}, p-value = {
441         p_value:.3f}")
442
443         if p_value < 0.05:
444             print("Conclusion: Algorithm B is significantly better than Algorithm
445             A.")
446         else:
447             print("Conclusion: No statistically significant difference between A
448             and B.")
449
450
451     def monte_carlo_simulation(num_trials: int = 100) -> tuple[np.ndarray, np.
452         ndarray]:
453         """
454         Run the Monte Carlo simulation to compare Algorithms A and B across random
455         scenarios.
456
457         Args:
458             num_trials: (int) The number of trials to run the simulation for.
459             Defaults to 50.
460
461         Returns:
462             tuple[ndarray[Any, float], ndarray[Any, float]]: results_a, results_b
463             """
464         a = simulate_committee_federated
465         b = simulate_proposed
466
467         results_a = zeros(num_trials)
468         results_b = zeros(num_trials)
469
470         for i in tqdm(range(num_trials), desc="simulation trials"):
471             alpha = sample_alpha()
472
473             results_a[i] = algorithm(a, **alpha)
474             results_b[i] = algorithm(b, **alpha)
475
476             hypothesis_test(results_a, results_b, kind="bootstrap")
477
478         return results_a, results_b
479
480
481     def visualize_algorithm_comparison(results_a, results_b, show_plots=True):
482         """
483         Visualize the comparison between two algorithms through various plots.

```

```

474
475     This function creates three visualization plots:
476     1. Histogram of scores for both algorithms
477     2. Histogram of differences between algorithm scores
478     3. Scatter plot of paired algorithm results
479
480     Args:
481         results_a (ndarray): Array of scores from Algorithm A
482         results_b (ndarray): Array of scores from Algorithm B
483         show_plots (bool): Whether to display the plots immediately. Default
484             is True.
485
486     Returns:
487         tuple: Three matplotlib figure objects (hist_fig, diff_fig,
488         scatter_fig)
489         """
490         # Calculate differences
491         differences = results_b - results_a # difference: B - A
492
493         # -----
494         # 1. Scatter Plot of Paired Data
495         # -----
496         scatter_fig = plt.figure(figsize=(8, 6))
497         plt.scatter(results_a, results_b, alpha=0.5, color='purple', edgecolor='
498         none')
499         plt.xlabel("Algorithm A Score")
500         plt.ylabel("Algorithm B Score")
501         plt.title("Scatter Plot: Algorithm B vs. Algorithm A")
502         plt.plot(
503             [min(results_a), max(results_a)], [min(results_a), max(results_a)],
504             linestyle='--', color='gray', label='Equality Line',
505             )
506         plt.legend()
507
508         # -----
509         # 2. Histogram each Score
510         # -----
511         hist_fig = plt.figure(figsize=(10, 6))
512         plt.hist(results_a, bins=20, label="Algorithm A", alpha=0.5)
513         plt.hist(results_b, bins=20, label="Algorithm B", alpha=0.5)
514         plt.xlabel("Fault Tolerance Score")
515         plt.ylabel("Frequency")
516         plt.title("Histogram of Fault Tolerance Score")
517         plt.legend()
518
519         # -----
520         # 3. Histogram of Differences
521         # -----
522         diff_fig = plt.figure(figsize=(8, 6))
523         sns.histplot(differences, bins=30, kde=True, color='skyblue', edgecolor='
524         black')
525         plt.xlabel("Difference (Algorithm B - Algorithm A)")
526         plt.ylabel("Frequency")
527         plt.title("Histogram of Differences in Scores")
528         plt.axvline(
529             np.mean(differences), color='red', linestyle='dashed', linewidth=2,
530             label=f'Mean Difference = {np.mean(differences):.2f}',
531             )
532         plt.legend()

```

```
529     if show_plots:  
530         plt.show()  
531  
532     return hist_fig, diff_fig, scatter_fig  
533  
534  
535 if __name__ == "__main__":  
536     # Run the simulation  
537     results_a, results_b = monte_carlo_simulation(num_trials=1000)  
538  
539     visualize_algorithm_comparison(results_a, results_b)
```

Listing 1: Python code for the Monte Carlo simulation experiment